Full Paper

Reading Without Having Read: How Reading Strategy, Critical Framework, and Textual Grounding Shape Computational Literary Discussion

Chris Brew[1,2]

The Ohio State University[1] LexisNexis Legal & Professional[2]
brew.2@osu.edu

A reader who surveys a novel broadly, attending to character and plot across many chapters, and a reader who studies it narrowly, concentrating on passages rich in social argument and thematic complexity, will produce recognisably different evidence bases for discussion. Yet when we implement these two reading strategies computationally and use them to generate multi-voice literary discussions about 15 novels, the resulting conversations are near-identical in character: expert airtime is stable to ±2% regardless of strategy. The two algorithms agree on fewer than 6% of their passage selections (Jaccard 0.057 across 60,774 passages), yet the interpretive framework—not the reading strategy—determines what gets said.

This convergence of synoptic and analytical reading provides large-scale computational evidence for Fish's thesis that meaning is reader-constructed. We test this finding alongside two further claims: that removing textual grounding produces not random hallucination but measurable pastiche*—graduated borrowing from the novel's own vocabulary, with fidelity ranging from 57% (Middlemarch) to 0% (Hester)—and that conversational scaffolding uniformly transforms monologue into dialogue across all 15 novels, independently of content. Together, these results decompose literary discussion into three independently measurable factors: reading strategy, interpretive framework, and grounding. Of these, framework dominates, grounding gates accuracy, and strategy is surprisingly irrelevant.*

1. Introduction: Two Ways of Reading, One Discussion

Consider two readers preparing to lead a seminar on *Middlemarch*. The first reads broadly, marking passages across thirty chapters that develop character and advance the plot—a *synoptic* reader. The second reads selectively, drilling into a smaller set of passages dense with social critique and thematic argument—an *analytical* reader. Both are legitimate scholarly strategies. They will arrive at the seminar with different evidence. The question is whether they will arrive at different discussions.

We can now answer this question computationally. We implement the synoptic and analytical strategies as two passage-selection algorithms applied to 15 Victorian and early-modern novels (60,774 passages), then generate structured multi-voice literary discussions from each selection. The synoptic method (min-cost network flow) selects 59 passages per run across 28 chapters, favouring character and plot. The analytical method (contextual embeddings) selects 32 passages from 20 chapters, favouring social critique and thematic depth. The two agree on fewer than 6% of their selections.

Yet the resulting discussions converge. Each expert's share of airtime varies by less than 2 percentage points regardless of which algorithm did the selecting. The interpretive framework—embodied in expert personas with distinct critical orientations—is the dominant factor shaping the discussion. The reading strategy is not.

This *framework dominance* is the central finding of this paper. It provides large-scale computational evidence for Stanley Fish's thesis that meaning is constituted by interpretive communities, not extracted from texts (Fish 1980). Different ways of reading the same novel produce different evidence bases but the same analytical character—because the reader, not the reading, makes the interpretation.

We build on this finding in three directions:

1. We show that the convergence holds across **15 novels** spanning seven authors and seven decades, with interpretive frameworks maintaining stable character regardless of the novel's affordances (Section 4).
2. We show that removing textual grounding produces not gibberish but **measurable pastiche**—the LLM recombines the novel's own vocabulary into plausible-but-invented quotations, with fidelity that tracks novel prominence from Middlemarch (57% verified) to Hester (0%). This provides an unintended instrument for measuring an LLM's literary knowledge (Section 6).
3. We show that conversational scaffolding transforms monologue into dialogue **uniformly across all 15 novels**, confirming that discussion dynamics are architecture-driven, not content-dependent (Section 7).

2. From Reading Strategies to Interpretive Frameworks

Our framework sits at the intersection of three traditions in literary theory.

Fish: the reader makes the meaning.. Fish (Fish 1980) argued that texts do not contain determinate meanings; meaning is constituted by the interpretive strategies readers bring. Two readers with different strategies will produce different readings of the same text—not because they disagree about the text but because they are, in effect, reading different texts. Our framework dominance finding extends this claim: not only different interpretive strategies but different *reading strategies*—synoptic vs. analytical, breadth vs. depth—produce convergent discussions, so long as the interpretive framework is held constant. It is not merely what you read but *who you are as a reader* that determines what you say.

Iser: the text affords, the reader fills.. Iser (Iser 1978) proposed that texts contain *Leerstellen*—gaps that readers fill according to their own resources. We operationalise affordance through seven provision dimensions (`character_development`, `plot_advancement`, `thematic_depth`, `social_critique`, `humor_entertainment`, `atmosphere_setting`, `narrative_technique`) applied to every passage in each novel. These dimensions vary across novels, confirming that different texts afford different interpretive possibilities—even as the framework dominance finding shows that interpretive frameworks override these affordances in determining what gets said.

Table 1
Corpus: 15 novels, 60,774 passages. Novels ordered by ungrounded verification rate (Section 6).

Novel	Author	Year	Passages	Nop%
Middlemarch	Eliot	1871	5,519	57
Bleak House	Dickens	1853	6,916	53
David Copperfield	Dickens	1850	5,621	44
Hard Times	Dickens	1854	1,773	39
A Passage to India	Forster	1924	2,556	39
Mill on the Floss	Eliot	1860	4,654	37
Cranford	Gaskell	1853	1,226	33
Daniel Deronda	Eliot	1876	5,170	30
Our Mutual Friend	Dickens	1865	5,775	29
New Grub Street	Gissing	1891	3,872	20
The Odd Women	Gissing	1893	3,395	6
Miss Marjoribanks	Oliphant	1866	3,019	3
North and South	Gaskell	1855	2,942	2
Hester	Oliphant	1883	2,279	0
No Name	Collins	1862	6,057	0

Moretti and Underwood: distant reading.. The distant reading programme (Moretti 2013; Underwood 2019) has analysed what novels *contain*. Our contribution is a different register: we analyse what novels *afford*—what kinds of discussions they enable when filtered through explicit critical lenses.

3. System and Experimental Design

We generate structured literary discussion podcasts through a pipeline that decomposes interpretation into manipulable phases. The system is the experimental instrument; we describe it briefly and refer readers to the companion paper for implementation details.

Corpus.. Fifteen novels (Table 1), parsed into passages averaging 50–80 words. Total: 60,774 passages across seven authors and seven decades (1850–1924).

Enrichment (Phase 0).. Each passage is annotated with a 20-field schema including an interest score (0–5), seven provision dimensions, characters present, themes, emotional register, and a best extractable quote.

Two reading strategies (Phase 1–2).. The synoptic and analytical reading strategies correspond to two passage-selection algorithms:

- **Synoptic (transport)**: min-cost network flow matching expert demand profiles to passage provisions. Selects broadly (59 passages, 28 chapters per run), favouring `character_development` (+8.3 pp vs. embedding) and `plot_advancement` (+10.2 pp). Zero LLM tokens, <1 second.
- **Analytical (embedding)**: contextual embeddings with LLM-curated similarity. Selects narrowly (32 passages, 20 chapters), favouring `social_critique` (+13.4 pp) and `thematic_depth` (+6.9 pp). ~15K Sonnet tokens, ~10 seconds.

Table 2
Synoptic vs. analytical reading: passage divergence and script convergence across 28 matched novel×panel pairs.

Measure	Synoptic	Analytical
Passages per run	59	32
Chapters per run	28	20
Mean interest score	4.44	4.01
`character_development` strong	83.5%	75.2%
`plot_advancement` strong	42.6%	32.5%
`thematic_depth` strong	84.0%	90.9%
`social_critique` strong	50.6%	64.0%
`humor_entertainment` strong	28.7%	24.4%
Passage Jaccard similarity	0.057 (range 0.00–0.12)	
Expert airtime SD across novels	<2 pp	

A third condition, **no passages**, removes grounding entirely: the LLM generates from prior knowledge alone.

Host preparation (Phase 2.5, optional).. Pre-interviews with each expert persona surface their views on assigned passages; question planning then creates targeted host questions.

Script generation (Phase 3).. One LLM call per segment produces a structured podcast script with speaker turns, quotations, and prosodic annotations.

Experimental design.. 180 conditions: 15 novels × 2 expert panels × 3 pipelines (synoptic, analytical, no-passages) × 2 host-preparation settings. Each condition produces one complete episode.

4. The Convergence Paradox

4.1 Different Passages, Same Discussion

Table 2 summarises the divergence between the two reading strategies and the convergence of the resulting discussions.

The two strategies select nearly disjoint passage sets (Jaccard 0.057), with systematically different provision profiles. The synoptic reader covers more of the novel, favouring passages strong in character and plot. The analytical reader concentrates on fewer passages, favouring those strong in social critique and thematic depth. These are recognisable scholarly strategies—the seminar leader who assigns the whole novel versus the one who curates a dossier of key passages.

Yet the discussions they produce are convergent. Expert airtime—each expert's share of total expert words, excluding the host—is stable across all 15 novels and both reading strategies: Eleanor Hartley (literary critic) 32–37%, James Blackstone (social historian) 30–37%, Caroline Woodcourt (close reader) 29–34%. The standard deviation across novels is less than 2 percentage points per expert.

4.2 Framework Stability, Not Adaptation

This stability was not our prediction. We originally hypothesised that expert prominence would *adapt* to textual affordance—that a novel rich in social critique would give the social historian more airtime. The data shows the opposite: frameworks are *robust*, maintaining their character regardless of what the novel provides. For digital humanities, this is the more interesting finding. The interpretive frameworks function as genuine critical lenses—stable heuristics for reading, not chameleons that reshape themselves to fit each text.

A caveat: framework dominance is realised through a pipeline that consistently privileges expert identity at every stage. Enrichment dimensions are defined in terms that match expert demand profiles; the passage selector optimises for demand satisfaction; host preparation pre-interviews each expert and plans questions targeting specific speakers; and the script-generation prompt explicitly instructs experts to maintain their perspectives and react to each other. Whether this finding reflects a deep property of interpretive discourse or the pipeline's engineering is an open question.

4.3 Fish's Thesis, Computationally

The framework dominance finding operationalises Fish's thesis at scale. Two recognisably different reading strategies produce convergent discussions because the dominant variable is the interpretive community (the expert panel), not the mode of reading. A Marxist survey reading and a Marxist close reading of *Hard Times* differ in their evidence base but not in their analytical character. A formalist reading of *Cranford* and a formalist reading of *Middlemarch* differ in their textual material but not in what the formalist notices.

The reader makes the reading. The reading strategy does not.

5. Textual Affordances as Novel Fingerprints

If interpretive frameworks dominate, do textual affordances matter at all? They do—but as *constraints on evidence*, not as determinants of interpretation.

Our seven provision dimensions produce distinctive profiles for each novel. The most diagnostic dimension is `humor_entertainment`, with a coefficient of variation of 0.704: *Miss Marjoribanks* (Oliphant) has the richest comedic supply (27.1% of passages strong), while *New Grub Street* (Gissing) and *The Odd Women* (Gissing) are among the sparsest. *Hard Times* affords social critique (20.2%) more than any other novel. *Miss Marjoribanks* affords character study (40.5%).

These affordance profiles are not subjective—they are reproducible computational measurements. They constitute a novel fingerprint: a measurable description not of what a novel *is* but of what it *affords* for discussion through specific critical lenses. This is Iser's (Iser 1978) concept of textual affordance, operationalised at scale.

Six of the seven dimensions transfer across all 15 novels. The exception, `atmosphere_setting`, is uniformly low (<6% strong), suggesting a bias in the schema toward features diagnostic in Realist fiction. The schema was designed on *Bleak House* and reflects Dickensian priorities. We regard this as a productive limitation: the perspective-dependence of descriptive schemas is itself a finding about literary description.

6. Reading Without Having Read

The framework dominance finding shows that reading strategy matters less than interpretive framework. But what if the reader has not read the novel at all?

We test this by removing passage grounding entirely. In the *no-passages* condition, the LLM generates literary discussion from prior knowledge alone—it must quote, analyse, and discuss without being given any text from the novel. This is the computational equivalent of the seminar participant who has not done the reading.

6.1 The Grounding Gap

When given passages (either strategy), the system quotes accurately: 98.0% verified (synoptic) and 97.2% (analytical). Without passages, the mean drops to 48.9%. This **grounding gap** ranges from ∼30 percentage points for canonical novels to ∼100 for obscure ones (Table 1, rightmost column).

6.2 The Knowledge Gradient

The ungrounded verification rate varies dramatically across novels: Middlemarch (57%), Bleak House (53%), David Copperfield (44%) at the top; Odd Women (6%), Miss Marjoribanks (3%), Hester and No Name (0%) at the bottom. The LLM can quote canonical novels from memory but cannot quote obscure ones at all.

This ordering does not simply track fame. Our Mutual Friend (Dickens, 29%) is more famous than Cranford (Gaskell, 33%). But the general pattern—declining fidelity with declining prominence—is clear.

6.3 Pastiche, Not Hallucination

The most distinctive finding concerns the *nature* of ungrounded quotation. We performed a vocabulary autopsy on all 726 invented quotes (0% match with any passage in the novel):

- **82%** draw 90–100% of their words from the novel's own vocabulary.
- The mean is **95.1%**—nearly all words appear somewhere in the novel.
- The minimum is 67%: even the worst inventions are built from the novel's word stock.

The LLM is writing *in the style of* the novel rather than *from* the novel. It recombines real vocabulary into plausible-but-nonexistent sequences. In literary-critical terms, this is *pastiche*—imitation of an author's manner without direct quotation.

For digital humanities, this reframes the "hallucination problem." Ungrounded LLM generation is not gibberish but creative imitation—sometimes convincing, always graduated, and novel-specific. The graduated nature makes it measurable: confabulation rate functions as an **unintended instrument for measuring LLM literary knowledge**. A novel the model can quote from memory is one it has deeply absorbed. A novel it cannot quote at all is one it has barely encountered.

7. Conversational Scaffolding

The framework dominance finding concerns the *content* of discussion. But literary discussions have *dynamics* too—the interplay of question, response, agreement, and disagreement that makes conversation conversational. Are these dynamics content-driven or can they be designed independently?

Host preparation—a phase in which the system pre-interviews each expert and plans targeted questions—transforms generated discussions. Without it, the system produces extended monologues ($\sim$1 question per segment). With it, the mean rises to $\sim$11 questions per segment.

This effect is **uniform across all 15 novels**: the ratio ranges from 5.3$\times$ (*Our Mutual Friend*) to 14.1$\times$ (*Cranford*), but every novel shows the transformation. The dynamics of literary discussion can be designed independently of its substance, through structural interventions at the level of conversation architecture.

This finding parallels the framework dominance finding. Just as reading strategy does not determine analytical character, literary content does not determine conversational dynamics. Both are shaped by the framework—interpretive in one case, architectural in the other.

8. Editorial Control: Exploring Interpretive Space

If the interpretive framework is the dominant factor, can it be manipulated?

The synoptic (transport) method formulates passage selection as min-cost flow in which each expert has a *demand profile* over the seven provision dimensions. By modifying these profiles—increasing one expert's appetite for social critique, another's for narrative technique—a scholar can explore how different emphases reshape the evidence base.

We tested graduated modifications, from single-dimension shifts to full panel replacement. Passage-set Jaccard values show a predictable response: conservative shifts produce 0.7–0.8 overlap with the baseline; radical replacement produces 0.5–0.6. Editorial intention is legible as a graduated flow of passage changes.

This exploration incurs **zero marginal LLM cost** (<100ms, no API calls), enabling a "preview-before-commit" workflow in which a scholar examines dozens of interpretive configurations before committing to expensive script generation. In Boden's (Boden 2004) terms, this is a tool for *exploratory creativity*—systematic traversal of interpretive space.

9. Discussion

9.1 A Decomposition of Literary Discussion

Our findings decompose computational literary discussion into three factors:

Interpretive framework determines the *character* of discussion: what the experts notice, how they frame it, what they argue about. This is the dominant factor. It is stable across novels and reading strategies.

Textual grounding determines *accuracy*: whether the discussion quotes the novel faithfully or produces pastiche. Without grounding, the framework still operates—the discussion sounds like literary criticism—but its evidentiary foundation collapses.

Reading strategy determines the *evidence base*: which passages are available for discussion. Surprisingly, this has minimal effect on the discussion's analytical character.

This decomposition has a clear hierarchy: framework > grounding > strategy. The finding that strategy is dominated by framework is the framework dominance finding. The finding that grounding gates accuracy but not fluency is the pastiche result.

9.2 The Schema Problem

Our enrichment schema transfers to 14 novels beyond *Bleak House*, with one dimension (`atmosphere_setting`) failing. The schema is inflected by its origin in Victorian Realism. A schema from modernist fiction would capture different affordances. This perspective-dependence is itself a finding: descriptive schemas, like interpretive frameworks, carry the priorities of their designers.

9.3 Limitations

Automated metrics only. No human evaluation of discussion quality, insight, or literary value.

English canon, 1850–1924. We cannot claim generalisability to other languages, periods, or genres.

Enrichment bias. The provision dimensions reflect Victorian Realist priorities.

No causal claims. Novel prominence correlates with verification rate, but we do not establish the causal mechanism.

Framework as prompt. Our "interpretive frameworks" are LLM personas specified in natural language. Whether they replicate the cognitive structures Fish describes is an open question.

10. Related Work

Computational literary criticism.. Moretti's (Moretti 2013) "distant reading" and Underwood's (Underwood 2019) statistical analyses of literary history have established computational methods in literary studies. Our work differs in generating *discourse about* literature rather than analysing literature directly, and in operationalising reader-response concepts (frameworks, affordances) rather than extracting textual features.

AI-generated conversation.. Google's NotebookLM (Google 2024) generates audio discussions from uploaded documents. Our contribution is the explicit decomposition into manipulable components, enabling the controlled comparisons that yield the framework dominance finding.

Retrieval-augmented generation.. Lewis et al. (Lewis et al. 2020) and Gao et al. (Gao et al. 2024) establish RAG as the standard approach to grounding LLM output. Our finding that the *method* of retrieval matters far less than the *fact* of retrieval extends this literature into a domain (literary discussion) where the relationship between source and output is unusually transparent.

Hallucination.. Our pastiche finding—that confabulation is structured, graduated, and vocabulary-preserving—characterises hallucination as a literary phenomenon rather than solely a technical failure.

11. Conclusion

A synoptic reader and an analytical reader, given the same novel and the same interpretive framework, produce discussions whose analytical character is determined by the framework, not the reading strategy. This convergence—demonstrated across 15 novels and 180 experimental conditions—provides the first large-scale computational evidence for Fish's thesis that meaning is reader-constructed.

When grounding is removed entirely, the system does not fall silent but produces graduated pastiche: plausible quotations constructed from the novel's own vocabulary, more faithful for canonical novels and degrading to pure invention for obscure ones. This pastiche provides a novel instrument for measuring LLM literary knowledge.

Conversational dynamics, like analytical character, are framework-driven: scaffolding transforms monologue into dialogue uniformly across all 15 novels. Editorial control through demand-profile manipulation enables zero-cost exploration of interpretive space.

These findings suggest that computational frameworks can contribute to—not replace—the humanities' understanding of how reading works. The reader makes the reading. The reading strategy does not.

References

Boden, Margaret A. 2004. *The Creative Mind: Myths and Mechanisms*, 2nd edition. Routledge.

Fish, Stanley. 1980. *Is There a Text in This Class? The Authority of Interpretive Communities*. Harvard University Press.

Gao, Yunfan, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. 2024. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*.

Google. 2024. NotebookLM now lets you listen to a conversation about your sources. `https://blog.google/technology/ai/notebooklm-audio-overviews/`.

Iser, Wolfgang. 1978. *The Act of Reading: A Theory of Aesthetic Response*. Johns Hopkins University Press.

Lewis, Patrick, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33.

Moretti, Franco. 2013. *Distant Reading*. Verso.

Underwood, Ted. 2019. *Distant Horizons: Digital Evidence and Literary Change*. University of Chicago Press.