

The Reader Makes the Reading: Framework Dominance and Pastiche in LLM-Generated Literary Discussion

Anonymous

Abstract

We present a system for generating multi-voice literary discussions about novels that decomposes interpretation into independently manipulable components: expert personas (critical frameworks), passage selection (reading strategies), and conversational scaffolding. Testing across 15 novels and 180 conditions, we find a *framework dominance*: two algorithmically distinct reading strategies—one synoptic, one analytical—select nearly disjoint passage sets (Jaccard 0.057) yet produce discussions whose character is determined by the interpretive framework, not the reading. This makes the system a tool for *exploratory creativity* in Boden’s sense: systematic traversal of an interpretive space defined by critical perspectives and textual affordances, where passage exploration is free and script generation costs only LLM tokens. Removing textual grounding reveals that ungrounded generation produces not random output but graduated *pastiche* that recombines the novel’s vocabulary—a form of combinational creativity whose fidelity tracks training-data exposure. We argue that the system’s value lies not in producing creative artefacts but in making the *space of possible interpretations* navigable.

Introduction

Computational creativity research has long distinguished between *exploratory* creativity—traversing a defined conceptual space—and *combinational* creativity—producing novel combinations of familiar elements (Boden 2004). Most LLM-based creative systems operate in combinational mode: they generate novel text by recombining learned patterns. We present a system that enables exploratory creativity in a specific domain: literary interpretation.

The system generates structured multi-voice podcast discussions about novels. Three expert personas, each embodying a distinct critical framework (e.g. formalist, social-historical, reader-response), discuss passages selected by an algorithmic reading strategy. The key insight is that the components are independently variable: the same novel can be discussed through different lenses, using different passage selections, with different conversational scaffolding—and each variation is a distinct point in an interpretive space that a human editor can explore.

This exploration is *cheap*. The passage-selection step uses min-cost network flow (Ahuja, Magnanti, and Orlin 1993) and runs in under 100ms with no LLM calls. An editor can preview dozens of interpretive configurations before committing to the expensive generation step. This “preview-before-commit” workflow distinguishes the system from generative tools like NotebookLM (Google 2024), which produce a single output with no mechanism for systematic variation. It realises the kind of “shared collaborative space” advocated by Ocampo, Bown, and Grace (2025), in which a human editor has genuine agency over the creative parameters rather than merely prompting a chat interface.

The Convergence Paradox

Our central empirical finding concerns the relative power of different creative levers. We implement two reading strategies as passage-selection algorithms applied to 15 Victorian novels (60,774 passages, 7 authors, 1850–1924):

- **Synoptic:** min-cost flow selects 59 passages from 28 chapters, favouring character development (+8.3 pp) and plot (+10.2 pp).
- **Analytical:** contextual embeddings select 32 passages from 20 chapters, favouring social critique (+13.4 pp) and thematic depth (+6.9 pp).

These are recognisably different scholarly strategies—the seminar leader who assigns the whole novel versus the one who curates a dossier of key passages. They agree on fewer than 6% of their selections (mean Jaccard 0.057 across 28 novel×panel pairs; see Table 1 for the full provision-dimension comparison).

Yet the resulting discussions converge: expert airtime is stable to $\pm 2\%$ across all 15 novels and both strategies (Hartley 32–37%, Blackstone 30–37%, Woodcourt 29–34%), regardless of which algorithm selected the passages or which novel is being discussed.

This stability was not predicted. We expected expert prominence to *adapt* to the novel’s affordances—that a novel rich in social critique would give the social historian more airtime. The data shows the opposite: frameworks are robust, maintaining their character across 15 novels spanning Dickens, Eliot, Gaskell, Forster, Collins, Gissing, and Oliphant.

Table 1: Provision profiles of selected passages (% rated ‘strong’), averaged across 28 novel×panel pairs.

Dimension	Synop.	Analyst.	Δ
character_development	83.5	75.2	+8.3
plot_advancement	42.6	32.5	+10.2
thematic_depth	84.0	90.9	−6.9
social_critique	50.6	64.0	−13.4
humor_entertainment	28.7	24.4	+4.3
narrative_technique	90.1	92.5	−2.4

A caveat is needed. Framework dominance is not a single mechanism but the cumulative effect of a pipeline that consistently privileges the expert’s identity at every stage: enrichment dimensions are defined in terms that match expert demand profiles; the passage selector optimises for expert-demand satisfaction; host preparation pre-interviews each expert individually and plans questions that target specific experts; and the script-generation prompt explicitly instructs speakers to maintain their distinctive perspectives and react to each other’s claims. Each stage reinforces the framework. Whether the finding reflects a deep property of interpretive discourse or merely the pipeline’s engineering is an open question—but the engineering is the mechanism, and we should name it honestly.

In Boden’s framework, this means the *conceptual space* of literary discussion is shaped primarily by the interpretive framework—the critical lens—not by the input material. The space is framework-defined, not text-defined. This is a structural finding about where creative agency resides in the system: it resides in the choice of perspective, not in the selection of evidence.

Pastiche as Combinational Creativity

What happens when textual grounding is removed entirely? In our *no-passages* condition, the LLM must discuss the novel from prior knowledge alone. The result is not silence or incoherence but graduated *pastiche*: the system constructs plausible quotations by recombining the novel’s own vocabulary. Of 726 invented quotes, 82% draw 90–100% of their individual words from the novel’s word stock. The mean overlap is 95.1%.

This is combinational creativity in a precise sense: the system produces novel combinations of familiar elements (the novel’s vocabulary) that are locally plausible but globally unfaithful.

The knowledge gradient. The fidelity of pastiche tracks cultural prominence: *Middlemarch* (57% verified), *Bleak House* (53%), *David Copperfield* (44%) at the top; *New Grub Street* (20%), *The Odd Women* (6%), *Hester* (0%) at the bottom. All 15 novels are on Project Gutenberg and plausibly in the LLM’s training data; the gradient reflects not textual presence but the volume of secondary discussion (study guides, academic papers, Wikipedia) that reinforces specific passages (Carlini et al. 2021; 2023). Bayard (2007) argued that humans navigate literary culture through graduated approximation; the LLM does something analogously measurable, though by entirely different mechanisms.

Table 2: The knowledge gradient: ungrounded quote verification (%) by novel. All novels are on Project Gutenberg; the gradient reflects analytical reinforcement, not textual absence. Grounded conditions achieve 97–98% for all novels.

Novel	Author	Nop%	Gap
Middlemarch	Eliot	57	42 pp
Bleak House	Dickens	53	45 pp
David Copperfield	Dickens	44	53 pp
Hard Times	Dickens	39	59 pp
Passage to India	Forster	39	55 pp
Mill on the Floss	Eliot	37	59 pp
Daniel Deronda	Eliot	30	69 pp
Our Mutual Friend	Dickens	29	69 pp
New Grub Street	Gissing	20	79 pp
The Odd Women	Gissing	6	91 pp
Hester	Oliphant	0	98 pp
No Name	Collins	0	95 pp

Within-author evidence. The gradient persists within authors, controlling for prose style. Among Dickens’s four novels: *Bleak House* (53%) > *David Copperfield* (44%) > *Hard Times* (39%) > *Our Mutual Friend* (29%). Among Eliot’s three: *Middlemarch* (57%) > *Mill on the Floss* (37%) > *Daniel Deronda* (30%). Same author, same prose style, different cultural prominence, different pastiche fidelity.

A concrete example illustrates the mechanism. The ungrounded system produces the quote “I am very weak, but I shall begin the world” and attributes it to *Bleak House*. Every word appears somewhere in the novel—100% vocabulary overlap—but this exact sentence does not exist. It is plausible Dickensian sentiment, constructed from the novel’s own word stock. Contrast this with the fog opening (“Fog everywhere. Fog up the river. . .”), which the system reproduces verbatim at 100% match. The boundary between faithful quotation and creative pastiche is measurable through our five-tier scheme (verified, paraphrase, distant echo, no clear source, invented).

For computational creativity, this finding matters because it characterises a *failure mode* as a *creative mode*. The system is not hallucinating randomly; it is performing pastiche—imitation of an author’s manner without direct quotation. It produces novel, plausible, stylistically coherent text, but creativity unmoored from its source.

Exploring Interpretive Space

The system’s practical value for computational creativity lies in its support for exploratory traversal of interpretive space. An editor can:

1. **Swap critical lenses:** replace a formalist with a Marxist, an astronomer, or an NLP scientist, and observe how the discussion reshifts.
2. **Adjust demand profiles:** increase one expert’s appetite for social critique or narrative technique, producing graduated passage-set changes (Jaccard 0.5–0.8 with baseline).
3. **Add or remove scaffolding:** host preparation raises ques-

tion density from ~ 1 to ~ 11 per segment, uniformly across all novels.

4. **Compare grounded and ungrounded:** observe the pastiche gradient as grounding is removed, measuring how much of the discussion’s substance depends on actual textual evidence.

Passage-level exploration incurs no LLM cost (the solver runs locally in $<100\text{ms}$). The expensive step—script generation, which costs $\sim 100\text{K}$ Sonnet tokens per episode—is deferred until the editor has identified a promising configuration. This is Boden’s *exploratory creativity* operationalised as a human-AI workflow: the human defines the space (through persona and demand choices), the machine traverses it (through optimisation and generation), and the human evaluates the results.

Relation to Computational Creativity

We see three connections to the CC literature:

Creativity resides in the framework, not the output.

Colton and Wiggins (2012) argue that attributing creativity requires understanding the *process*, not just evaluating the *product*. Morain and Ventura (2025) ask whether prompt engineering is the “creativity knob” for LLMs. Our framework dominance finding provides an empirical answer: the persona prompt—not the textual input or retrieval method—determines the character of the output. The creative act is the choice of lens.

Pastiche is a measurable creative mode. Ritchie (2007) proposes criteria for attributing creativity to programs, including novelty and value. Pastiche is novel (the sequences do not exist in the source text) and locally valuable (they are stylistically coherent). But it lacks *fidelity*—it does not faithfully represent the source. Our five-tier verification scheme (verified, paraphrase, distant echo, no clear source, invented) provides a graded measure of this trade-off.

The system enables human exploratory creativity. Ismayilzada, Stevenson, and van der Plas (2025) find that LLMs generate stylistically complex stories but fall short in novelty and diversity compared to humans. Doshi and Hauser (2024) find that generative AI enhances individual creativity but reduces collective diversity. Our system addresses both concerns: by making the interpretive framework an explicit, manipulable parameter, it produces *diverse* outputs from the same source material. The system does not converge on a single “best” interpretation; it maps the space of possible interpretations.

What If I Put My Friends in the Podcast?

The system’s personas are not limited to literary scholars. To test the boundaries of exploratory creativity, we replaced the default panel (a formalist, a social historian, and a close reader) with three American professionals from outside literary studies: a computer scientist specialising in speech recognition and NLP, an observational astronomer who studies stellar formation, and a musicologist who studies Cold War cultural diplomacy.

The resulting discussions of *Bleak House* are recognisably different from the literary-critical baseline. The astronomer

reads Dickens’s fog as a physical formation condition—opacity is not metaphor but the baseline state of complex systems. The NLP scientist analyses agent deletion in the opening sentences: “There is no main clause. It is a sequence of noun phrases and participial clauses with no finite verb until paragraph two.” The musicologist compares Chancery to Soviet institutions that consumed the artists they were supposed to support.

Two observations matter for computational creativity:

First, the *framework dominance finding still holds*. Despite radically different personas, expert airtime remains balanced and the discussions are structurally similar to the literary-critical ones. The framework dominates even when the framework is non-literary.

Second, the experts bring their *metaphors* but not their *methods*. The astronomer does not propose measurements or hypotheses; she produces literary-critical discourse decorated with astronomical imagery. This is a limitation of the prompt architecture: the host preparation questions (“what strikes you about these passages?”) import literary-critical norms regardless of the persona. Sood, Llano, and McCormack (2025) ask whether conversational interfaces limit creativity; our finding suggests that conversational *scaffolding* can constrain disciplinary expression even when the persona is unconstrained. A prompt designed for genuinely interdisciplinary encounter—“what would your discipline’s tools reveal about this passage?”—might produce different results. This is itself a finding about where creative constraint operates: not just in the persona but in the conversational scaffold.

Debate Spark: Is Systematic Variation Creative?

The framework dominance finding makes one thing obvious: the personas and framework bear the creative load in this system—perhaps badly. The LLM is a *realiser*. It takes a specification (persona, passages, scaffolding instructions) and renders it into fluent multi-voice discourse. The rendering is impressive. But the creative decisions—which lens to apply, which dimensions to weight, which experts to put in the room together—are made by the human editor before the LLM is invoked.

This is not a controversial claim. It is the direct implication of the framework dominance finding: since the interpretive framework determines the output’s character regardless of reading strategy, the framework is where the creative agency sits, and the framework is human-authored.

What *is* debatable is whether the system’s ability to **vary systematically at scale** constitutes a creative contribution that deserves the CC community’s attention. We think it does, for three reasons:

1. Scale makes comparison possible. A human editor cannot commission 180 literary discussions to compare how different lenses reveal different aspects of 15 novels. The system can, and the comparison itself—the ability to see that a Marxist reading and a formalist reading of *Hard Times* select completely different passages but produce structurally similar discussions—is a finding that no single discussion

could produce. Creativity here lies not in any individual output but in the *landscape of outputs*.

2. Zero-cost exploration enables serendipity. Because passage selection runs in under 100ms with no LLM calls, an editor can try configurations they would never commission at full cost. “What if I put an astronomer in the panel?” is a question you ask when asking is free. The interdisciplinary panel (Section 7) emerged from exactly this dynamic: a whimsical configuration that produced genuinely surprising readings.

3. The pastiche gradient reveals something new. The graduated degradation from faithful quotation (57% for *Middlemarch*) to pure invention (0% for *Hester*) is not something the system was designed to produce. It is an emergent property of removing grounding from an LLM that has differentially memorised its training data. Whether this counts as “creativity” is precisely the kind of boundary question CC research should address.

Objections we take seriously:

The strongest objection is that systematic variation is *engineering*, not creativity. A tool that renders many versions of the same specification is an industrial process, not an artistic one. We would accept this objection if the variations were predictable—but they are not. The framework dominance finding was surprising. The pastiche gradient was surprising. The interdisciplinary panel’s “metaphors but not methods” finding was surprising. The system produces surprises not through generative novelty but through *systematic comparison*—and we argue that surprise-through-comparison deserves a place in the CC vocabulary alongside surprise-through-generation.

A second objection is that calling the human the “creative agent” diminishes the LLM’s contribution. The LLM produces the actual prose—hundreds of thousands of words of fluent, contextually appropriate literary discussion. Is that not creative? We think not, for the reason the framework dominance finding demonstrates: the prose is *determined* by the specification. Change the persona and the prose changes predictably. A system whose outputs are specification-determined is a realiser, however impressive its realisations. But we recognise this is a contested claim and invite the community’s response.

Demo

An interactive demo is available at <https://bleakhouse-demo.fly.dev>, where readers can browse the 180-condition experiment matrix, listen to generated audio episodes, and inspect passage backlinks, host preparation data, and quote verification results for all runs including the interdisciplinary panel.

Limitations

Our “interpretive frameworks” are LLM persona prompts, not cognitive models. As Shanahan (2024) warns, we should not map LLM behaviour onto human cognitive categories without justification. The personas are stereotypes, not theories—useful experimental instruments, not models of

creative cognition. We report automated metrics only, with no human evaluation of creative quality.

Author Contributions

Anonymised for review.

References

- Ahuja, R. K.; Magnanti, T. L.; and Orlin, J. B. 1993. *Network Flows: Theory, Algorithms, and Applications*. Prentice-Hall.
- Bayard, P. 2007. *How to Talk About Books You Haven’t Read*. Granta.
- Boden, M. A. 2004. *The Creative Mind: Myths and Mechanisms*. Routledge, 2nd edition.
- Carlini, N.; Tramèr, F.; Wallace, E.; Jagielski, M.; Herbert-Voss, A.; Lee, K.; Roberts, A.; Brown, T.; Song, D.; Erlings-son, Ú.; Oprea, A.; and Raffel, C. 2021. Extracting training data from large language models. In *USENIX Security Symposium*.
- Carlini, N.; Ippolito, D.; Jagielski, M.; Lee, K.; Tramèr, F.; and Zhang, C. 2023. Quantifying memorization across neural language models. *arXiv preprint arXiv:2202.07646*.
- Colton, S., and Wiggins, G. A. 2012. Computational creativity: The final frontier? *Frontiers in Artificial Intelligence and Applications* 242:21–26.
- Doshi, A. R., and Hauser, O. P. 2024. Generative AI enhances individual creativity but reduces the collective diversity of novel content. *Science Advances* 10(28).
- Google. 2024. NotebookLM now lets you listen to a conversation about your sources. blog.google/technology/ai/notebooklm-audio-overviews/.
- Ismayilzada, M.; Stevenson, C.; and van der Plas, L. 2025. Evaluating creative short story generation in humans and large language models. In *Proceedings of the International Conference on Computational Creativity*.
- Morain, R., and Ventura, D. 2025. Is prompt engineering the creativity knob for large language models? In *Proceedings of the International Conference on Computational Creativity*.
- Ocampo, R.; Bown, O.; and Grace, K. 2025. Beyond chat: Towards greater user involvement and agency in human-AI co-creativity through shared collaborative spaces. In *Proceedings of the International Conference on Computational Creativity*.
- Ritchie, G. 2007. Some empirical criteria for attributing creativity to a computer program. In *Minds and Machines*, volume 17, 67–99.
- Shanahan, M. 2024. Talking about large language models. *Communications of the ACM* 67(2):68–79.
- Sood, A.; Llano, M. T.; and McCormack, J. 2025. Do conversational interfaces limit creativity? Exploring visual graph systems for creative writing. In *Proceedings of the International Conference on Computational Creativity*.

Example Outputs: Literary vs. Interdisciplinary

The following excerpts show the opening turns from the same novel (*Bleak House*) discussed by two different panels under the same conditions (transport selection, host preparation).

Literary Panel (Hartley / Blackstone / Woodcourt)

[Caroline Woodcourt] I'll stand by it, though I'd refine it slightly. The fog is magnificent—I don't want to seem ungrateful for the fog. But readers sometimes use the fog as a kind of permission—'this is a great symbolic novel, I'm in the presence of Literature with a capital L'—and then they stop actually listening to what Dickens is telling them. What moves me in these passages is something far quieter—it's the moment where a human life is described as leaving no trace.

[James Blackstone] I take the point, but I'd push back—because the fog and the individual tragedy aren't separable, that's precisely what Dickens is insisting on. The fog is not decoration. It is the Court of Chancery made visible. "Lies there with no more track behind him that any one can trace than a deserted infant." That sentence is the whole novel compressed into a single image.

[Eleanor Hartley] What strikes me about that Nemo sentence is the grammatical audacity of it. 'A deserted infant'—Dickens reaches for the most helpless possible image of erasure. The architecture of that is devastating—the infant image should be about beginning, and Dickens flips it into a closing.

Interdisciplinary Panel (Chen / Martinez / Volkov)

[Sarah Chen] This is from chapter nineteen—the long vacation passage. "The good ships Law and Equity, those teak-built, copper-bottomed, iron-fastened, brazen-faced, and not by any means fast-sailing clippers." That is a noun phrase with four pre-nominal modifiers, each compounded—and then 'brazen-faced' pulls the floor out. The first three feel like nautical specifications—they're doing the register of technical competence—and the fourth is an insult.

[Elena Volkov] That structural camouflage is exactly what makes it so insidious. Institutions don't announce their corruption—they announce their solidity, their procedure, their continuity. I spent years studying how Soviet conservatories presented themselves as nurturing talent while systematically destroying composers who stepped out of line. The aesthetic of solidity covered everything.

[Rebecca Martinez] The ships of Law and Equity aren't sunk, they're in reserve. They'll be refloated. In my field, we'd say the system is in a quasi-stable state. The energy is there. Nothing has been resolved. The fog doesn't obscure something that would otherwise be clear. It is the natural atmosphere of a system that has never actually resolved anything.

Both panels produce engaged, multi-voice discussion. Woodcourt reads for emotional weight; Blackstone for institutional critique; Hartley for structural craft. Chen reads for

actual syntactic structure; Volkov for political parallel; Martinez for physical-systems analogy. The interpretive *character* is different; the conversational *dynamics* are convergent.

Vocabulary Autopsy: Measuring Pastiche

Our five-tier verification scheme classifies each quote utterance in ungrounded (no-passages) runs by matching against the full enriched passage corpus:

- **Verified** ($\geq 60\%$ 5-word-window match): genuine quotation from the novel.
- **Paraphrase** (30–59%): recognisable but altered.
- **Distant echo** (10–29%): loosest thematic connection.
- **No clear source** (1–9%): nearest passage shown.
- **Invented** (0%): no 5-word window matches any passage.

For invented quotes, a **vocabulary autopsy** reports how many of the quote's individual words appear anywhere in the novel. Examples from the interdisciplinary panel's ungrounded *Bleak House* discussion:

Verified (100%): "Fog everywhere. Fog up the river, where it flows among green aits and meadows; fog down the river, where it rolls defiled among the tiers of shipping." —*The famous opening; reproduced verbatim from memory.*

Invented (100% words in novel): "I am very weak, but I shall begin the world." —*Every word appears in Bleak House, but this exact sequence does not. Classic pastiche: plausible Dickensian sentiment constructed from the novel's own vocabulary.*

Invented (93% words in novel): "Both the long-locked up proprietors, each claiming the right to let the other go to ruin." —*One word ('proprietors') is absent from the novel. The rest is Dickensian pastiche of legal satire.*

Invented (80% words in novel): "I don't know nothink." —*Jo's characteristic speech; the LLM remembers the dialect but reconstructs the phrasing rather than reproducing a specific passage.*

The pattern is consistent across all 726 invented quotes: 82% draw 90–100% of their words from the novel. The LLM writes *in the style of* the text, not *from* it.

Why One Is Synoptic and the Other Analytical

The synoptic/analytical distinction is not an accident of parameter tuning. It follows from the mathematical structure of each method.

Transport (synoptic). The min-cost flow formulation (Ahuja, Magnanti, and Orlin 1993) models passage selection as an optimisation problem. Each passage is a supply node with provisions across seven dimensions (character development, plot advancement, social critique, etc.). Each expert is a demand node with an appetite for specific dimensions. The solver finds the minimum-cost assignment that satisfies all demands. Default budget: up to 60 passages, 8–25 per expert.

Because demand is *multi-dimensional*—each expert wants passages strong in several different dimensions—the

solver must draw from across the novel to satisfy all demands simultaneously. A single chapter rarely supplies all seven dimensions at once. The result is broad coverage: 59 passages from 28 chapters on average. The synoptic character is a consequence of multi-dimensional demand meeting a cost-minimising solver.

Strong matches (provision = “strong”) cost 1; weak matches cost 3. The solver prefers passages that strongly match an expert’s primary demand, which tends to select the most topically concentrated passages in each dimension—but across many dimensions, producing breadth.

Embedding (analytical). The embedding method retrieves passage candidates by semantic similarity to each expert’s profile description, then uses an LLM (Sonnet) to curate the final set. Default: 100 candidates retrieved, curated to 32 final passages.

Because retrieval is *similarity-driven*, passages cluster around each expert’s dominant theme. A social historian’s profile description embeds close to passages about institutional critique; retrieved candidates are topically concentrated. LLM curation selects the most relevant from this already-concentrated pool. The result is depth: 32 passages from 20 chapters, concentrated in the novel’s most thematically dense sections.

The analytical character is a consequence of cosine similarity producing topical clustering, compounded by LLM curation that selects for argumentative density over narrative breadth.

Is this systematic? Yes, but with a caveat. The passage counts (60 budget vs. 32 target) are design parameters, not mathematical necessities. Setting the transport budget to 32 would produce narrower coverage; setting the embedding target to 60 would force broader retrieval. But the *qualitative* character of each method—multi-dimensional optimisation producing breadth, similarity-based retrieval producing depth—would persist at any matched budget. Transport distributes across dimensions; embedding concentrates within themes. This is an inherent property of the two algorithmic approaches, not a tuning artefact.

Convergence Summary

The passage Jaccard similarity of 0.057 (range 0.000–0.118) means the two strategies share on average 3–7 passages out of 50–60 selected. The provision-dimension profiles (Table 1 in the main text) confirm that the difference is systematic: synoptic favours character and plot; analytical favours social critique and thematic depth. Despite this near-disjointness, expert airtime is stable to ± 2 pp across all 15 novels—the interpretive lens, not the reading strategy, determines the character of the output.