READING WITHOUT HAVING READ: HOW READING STRATEGY, INTERPRETIVE FRAMEWORK, AND TEXTUAL GROUNDING SHAPE COMPUTATIONAL LITERARY DISCUSSION

Chris BREW[1, 2]

[1]The Ohio State University, Columbus, Ohio, USA

[2]LexisNexis Legal & Professional

A reader who surveys a novel broadly and a reader who studies it narrowly will assemble different evidence for discussion. Yet when we implement these two reading strategies computationally and generate multi-voice literary discussions about 15 novels (60,774 passages, 180 experimental conditions), the resulting conversations are near-identical in character: expert airtime is stable to ±2% regardless of strategy. The two algorithms agree on fewer than 6% of their selections, yet the interpretive framework—not the reading strategy—determines what gets said. This convergence provides computational evidence for Fish (1980)'s thesis that meaning is reader-constructed. Removing textual grounding produces not random hallucination but *measurable pastiche*: 82% of invented quotes draw 90–100% of their vocabulary from the novel's word stock, with fidelity ranging from 57% (*Middlemarch*) to 0% (*Hester*). This extends Bayard (2007)'s philosophical argument about non-reading into an empirically testable framework—and connects it to recent work on LLM memorisation of training data (Carlini et al., 2021, 2023).

Keywords: literary interpretation, reader-response theory, LLM memorisation, passage grounding, digital humanities

1 TWO WAYS OF READING, ONE DISCUSSION

Consider two scholars preparing to lead a seminar on *Middlemarch*. The first reads broadly, marking passages across thirty chapters that develop character and advance the plot—a *synoptic* reader. The second reads selectively, concentrating on passages dense with social critique and thematic argument—an *analytical* reader. Both are legitimate strategies. They will arrive at the seminar with different evidence. The question is whether they will arrive at different discussions.

We can now answer this question computationally. We implement the two strategies as passage-selection algorithms applied to 15 Victorian and early-modern novels, then generate structured multi-voice literary discussions from each selection. The synoptic method (min-cost network flow; Ahuja et al., 1993) selects 59 passages per run from 28 chapters, favouring character development and plot advancement. The analytical method (contextual embeddings with LLM-curated similarity) selects 32 passages from 20 chapters, favouring social critique and thematic depth. The two agree on fewer than 6% of their selections (Jaccard 0.057).

Yet the resulting discussions converge. Each expert's share of airtime varies by less than 2 percentage points regardless of which algorithm did the selecting. The interpretive framework—embodied in expert personas with distinct critical orientations—dominates. The reading strategy does not.

This *convergence paradox* is the central finding of this paper. It provides large-scale computational evidence for Fish (1980)'s thesis that meaning is constituted by interpretive communities, not extracted from texts. We build on it in three directions:

1. We show that convergence holds across 15 novels, with interpretive frameworks maintaining stable character regardless of textual affordance (Section 4).

2. We show that removing grounding produces *measurable pastiche*—extending Bayard (2007)'s philosophical argument about non-reading into an empirically testable framework, and connecting it to research on LLM memorisation (Carlini et al., 2021) (Section 5).

3. We show that conversational scaffolding transforms monologue into dialogue uniformly across all novels, independently of content (Section 6).

2 FROM READER-RESPONSE THEORY TO COMPUTATIONAL EXPERIMENT

2.1 Interpretive communities

Fish (1980) argued that texts do not contain determinate meanings; meaning is constituted by the interpretive strategies readers bring. Our convergence paradox extends this: not only different interpretive strategies but different *reading strategies*—synoptic vs. analytical—produce convergent discussions, so long as the interpretive framework is held constant.

2.2 Textual affordance

Iser (1978) proposed that texts contain gaps that readers fill according to their resources. We operationalise affordance through seven provision dimensions applied to every passage. These dimensions vary across novels, confirming that different texts afford different interpretive possibilities—even as the convergence paradox shows that frameworks override affordances.

2.3 Non-reading as a spectrum

Bayard (2007) argued provocatively that "non-reading" is not a binary state but a spectrum: books skimmed, books heard about, books forgotten. He suggested that cultural positioning matters more than textual knowledge. Our no-passages condition—in which the LLM generates literary discussion from prior knowledge alone—provides the first computational test of this claim. The results confirm Bayard's intuition, but with a crucial empirical twist: we can *measure* the graduated degradation from reading to non-reading, and show that it tracks the model's training-data exposure to each novel.

2.4 LLM memorisation and literary knowledge

Carlini et al. (2021) demonstrated that language models memorise and can reproduce verbatim excerpts from training data. Carlini et al. (2023) showed that memorisation grows log-linearly with model capacity, data duplication, and prompt length. Duarte et al. (2024) developed methods to detect which books are in a model's training data. Our pastiche finding connects this technical literature to literary-critical concerns: the LLM's differential ability to quote different novels is a measurable trace of its training-data composition, visible as a gradient from faithful quotation of canonical works to pure pastiche of obscure ones.

3 SYSTEM AND EXPERIMENTAL DESIGN

We generate structured literary discussion podcasts through a pipeline that decomposes interpretation into manipulable phases. The system is the experimental instrument; a companion technical paper describes implementation details.

Corpus. Fifteen novels (Table 1), parsed into passages of 50–80 words. Total: 60,774 passages across seven authors (Dickens, Eliot, Gaskell, Forster, Collins, Gissing, Oliphant) and seven decades (1850–1924).

Table 1: Corpus: 15 novels, 60,774 passages, ordered by ungrounded verification rate (Section 5).

Novel	*Author*	*Year*	*Passages*	*Nop%*
Middlemarch	Eliot	1871	5,519	57
Bleak House	Dickens	1853	6,916	53
David Copperfield	Dickens	1850	5,621	44
Hard Times	Dickens	1854	1,773	39
A Passage to India	Forster	1924	2,556	39
Mill on the Floss	Eliot	1860	4,654	37
Cranford	Gaskell	1853	1,226	33
Daniel Deronda	Eliot	1876	5,170	30
Our Mutual Friend	Dickens	1865	5,775	29
New Grub Street	Gissing	1891	3,872	20
The Odd Women	Gissing	1893	3,395	6
Miss Marjoribanks	Oliphant	1866	3,019	3
North and South	Gaskell	1855	2,942	2
Hester	Oliphant	1883	2,279	0
No Name	Collins	1862	6,057	0

Two reading strategies. The synoptic strategy (min-cost flow) selects 59 passages from 28 chapters, favouring `character_development` (+8.3 pp) and `plot_advancement` (+10.2 pp). The analytical strategy (embeddings) selects 32 from 20 chapters, favouring `social_critique` (+13.4 pp) and `thematic_depth` (+6.9 pp).

Experimental matrix. 180 conditions: 15 novels × 2 expert panels × 3 pipelines (synoptic, analytical, no-passages) × 2 host-preparation settings (with/without).

4 THE CONVERGENCE PARADOX

The two strategies select nearly disjoint passages with systematically different provision profiles (Table 2). The synoptic reader covers the novel broadly; the analytical reader concentrates on its arguments.

Yet expert airtime—each expert's share of total words, excluding the host—is stable across all 15 novels and both strategies: Eleanor Hartley (literary critic) 32–37%, James Blackstone (social historian) 30–37%, Caroline Woodcourt (close reader) 29–34%.

Table 2: Synoptic vs. analytical reading across 28 matched pairs.

Measure	*Synoptic*	*Analytical*
Passages per run	59	32
Chapters per run	28	20
Mean interest score	4.44	4.01
`character_development` strong	83.5%	75.2%
`plot_advancement` strong	42.6%	32.5%
`social_critique` strong	50.6%	64.0%
`thematic_depth` strong	84.0%	90.9%
Passage Jaccard	0.057 (range 0.00–0.12)	
Expert airtime SD	<2 pp across 15 novels	

This stability was not our prediction. We expected expert prominence to *adapt* to textual affordance. Instead, the frameworks are *robust*—stable heuristics for reading, not chameleons. In Fish (1980)'s terms: the reader makes the reading, and the reading strategy does not.

5 READING WITHOUT HAVING READ

5.1 The grounding gap

With passages, the system quotes accurately: 98.0% (synoptic), 97.2% (analytical). Without passages, the mean drops to 48.9%, ranging from 57% (*Middlemarch*) to 0% (*Hester*, *No Name*). This grounding gap ranges from ∼30 pp for canonical novels to ∼100 pp for obscure ones.

5.2 Bayard's spectrum, measured

Bayard (2007) argued that we always talk about "approximate recollections of books, arranged as a function of our current circumstances." The no-passages condition makes this literal: the LLM's "recollection" of each novel is whatever it absorbed during training. Our verification rates measure this recollection on a per-novel basis, producing a *knowledge gradient* that tracks—but does not perfectly follow—novel prominence.

Middlemarch (57% verified) and *Bleak House* (53%) are canonical texts whose passages the model can often reproduce from memory. *Hester* (Oliphant, 1883) and *No Name* (both 0%) are texts the model has barely encountered. The gradient between these extremes is itself a measurement of literary-cultural prominence as reflected in training data.

5.3 Pastiche, not hallucination

We performed a vocabulary autopsy on 726 invented quotes (0% match with any passage):

- 82% draw 90–100% of their words from the novel's vocabulary;
- mean overlap is 95.1%; minimum is 67%.

The LLM writes *in the style of* the novel, not *from* it. It recombines real vocabulary into plausible-but-nonexistent sequences. In literary terms, this is *pastiche*.

This connects to Carlini et al. (2021)'s finding that LLMs memorise training data differentially. Carlini et al. (2023) showed memorisation grows with data duplication—and canonical novels are duplicated far more than obscure ones in web-scraped training corpora. Our pastiche gradient is the literary-critical face of this technical phenomenon: the model's ability to quote *Middlemarch* reflects the frequency with which *Middlemarch* appeared in its training data.

Bender et al. (2021) warned that language models are "stochastic parrots" producing text without understanding. Our pastiche results give this metaphor empirical texture: the parrot's fidelity depends on how many times it has heard each text.

6 CONVERSATIONAL SCAFFOLDING

Host preparation (pre-interviews + question planning) raises questions per segment from $\sim$1 to $\sim$11, **uniformly across all 15 novels**. The ratio ranges from 5.3$\times$ (*Our Mutual Friend*) to 14.1$\times$ (*Cranford*). Just as reading strategy does not determine analytical character, literary content does not determine conversational dynamics. Both are shaped by framework—interpretive in one case, architectural in the other.

7 WEAKNESSES AND HONEST LIMITATIONS

We owe the reader a frank accounting of what this framework does *not* demonstrate, alongside what it does.

7.1 LLM prompts are not a theory of interpretation

Our “interpretive frameworks” are natural-language persona descriptions fed to an LLM. They produce outputs that *resemble* literary criticism performed from specific perspectives. But resemblance is not identity. As Shanahan (2024) argues, we should resist the temptation to map LLM behaviour onto human cognitive categories without justification. When we say the system “performs a Marxist reading,” we mean that its output contains vocabulary, framing, and emphasis characteristic of Marxist criticism—not that it executes the intellectual operations a Marxist critic would. The convergence paradox shows that persona prompts produce stable, distinctive outputs. It does *not* show that these outputs are produced by the same mechanisms as human literary interpretation.

To put it plainly: our personas are stereotypes, not theories. They are useful experimental instruments—they produce controlled, reproducible variation in a generated text—but they should not be mistaken for cognitive models of how human readers interpret literature. Fish (1980)’s interpretive communities are social and cognitive phenomena; our expert personas are prompt engineering. The convergence between them—that both produce framework-dominant, text-secondary interpretive behaviour—is suggestive but not probative.

7.2 Automated metrics only

We report verification rates, airtime distributions, and structural statistics, but no human evaluation of discussion quality, insight, or literary value. A discussion that scores well on our metrics could still be intellectually shallow. Human evaluation is the most important missing piece.

7.3 English canon, 1850–1924

All 15 novels are English, from a narrow period. We cannot claim generalisability to other languages, periods, or genres. The provision dimensions reflect Victorian Realist priorities (`atmosphere_setting` fails to transfer), and a schema developed from different traditions would capture different affordances.

7.4 No causal claims

Novel prominence correlates with verification rate, but we do not establish the causal mechanism. The correlation could reflect training-data duplication (Carlini et al., 2023), prose distinctiveness, or quotability—we cannot distinguish these.

8 RELATED WORK

Moretti (2013) and Underwood (2019) established computational methods in literary studies through topic modelling, network analysis, and stylometry. Our work differs in generating *discourse about* literature and in operationalising reader-response concepts.

Google's NotebookLM (Google, 2024) generates audio discussions from documents. Our contribution is the explicit decomposition into manipulable components, enabling the controlled comparisons that yield the convergence paradox.

Lewis et al. (2020) and Gao et al. (2024) establish retrieval-augmented generation as the standard approach to grounding. Our finding that retrieval *method* matters far less than the *fact* of retrieval extends this into a domain where source–output relationships are transparent.

9 CONCLUSION

A synoptic reader and an analytical reader, given the same novel and the same interpretive framework, produce discussions whose character is determined by the framework, not the reading strategy. This convergence—across 15 novels and 180 conditions—provides computational evidence, though not proof, for Fish (1980)'s thesis that meaning is reader-constructed.

When grounding is removed, the system produces graduated pastiche: plausible quotations built from the novel's vocabulary, more faithful for canonical novels and degrading to invention for obscure ones. This pastiche is the literary-critical face of differential LLM memorisation (Carlini et al., 2021), and gives empirical texture to Bayard (2007)'s argument that we are always talking about "approximate recollections" of books.

We close with a caution. These findings are suggestive, not dispositive. LLM personas are not theories of mind. Automated metrics are not literary judgements. But the convergence paradox, the pastiche gradient, and the uniformity of scaffolding effects are empirical regularities that any account of computational literary discussion must accommodate. The reader makes the reading. The reading strategy does not.

ACKNOWLEDGEMENTS

The author thanks the anonymous reviewers for their comments.

REFERENCES

Ahuja, R. K., Magnanti, T. L., & Orlin, J. B. (1993). *Network flows: Theory, algorithms, and applications*. Prentice-Hall.

Bayard, P. (2007). *How to talk about books you haven't read* [Originally published as *Comment parler des livres que l'on n'a pas lus?*, Éditions de Minuit]. Granta.

Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? *Proceedings of FAccT*.

Carlini, N., Ippolito, D., Jagielski, M., Lee, K., Tramèr, F., & Zhang, C. (2023). Quantifying memorization across neural language models. *arXiv preprint arXiv:2202.07646*.

Carlini, N., Tramèr, F., Wallace, E., Jagielski, M., Herbert-Voss, A., Lee, K., Roberts, A., Brown, T., Song, D., Erlingsson, Ú., Oprea, A., & Raffel, C. (2021). Extracting training data from large language models. *USENIX Security Symposium*.

Duarte, A. V., Fonseca, E., Kela, M., Klakow, D., & Semedo, D. (2024). DE-COP: Detecting copyrighted content in language models training data. *International Conference on Learning Representations*.

Fish, S. (1980). *Is there a text in this class? The authority of interpretive communities*. Harvard University Press.

Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., Wang, M., & Wang, H. (2024). Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*.

Google. (2024). NotebookLM now lets you listen to a conversation about your sources. https://blog.google/technology/ai/notebooklm-audio-overviews/

Iser, W. (1978). *The act of reading: A theory of aesthetic response*. Johns Hopkins University Press.

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, *33*.

Moretti, F. (2013). *Distant reading*. Verso.

Oliphant, M. (1883). *Hester*. Macmillan; Co.

Shanahan, M. (2024). Talking about large language models. *Communications of the ACM*, *67*(2), 68–79.

Underwood, T. (2019). *Distant horizons: Digital evidence and literary change*. University of Chicago Press.

APPENDIX: WHY CAN THE LLM QUOTE MIDDLEMARCH BUT NOT HESTER?

When asked to discuss a novel without source passages, the LLM produces verifiable quotations at rates ranging from 57% (*Middlemarch*) to 0% (*Hester*). All 15 novels are available from Project Gutenberg[1] and plausibly appear in the model's training data (though we cannot verify this for any proprietary model). If so, the gradient must reflect something other than raw textual presence. This appendix investigates what that something is, subjects the investigation to systematic critique, and states what survives.

A.1 Hypothesis: Analytical Reinforcement

Carlini et al. (2023) showed that LLM memorisation grows log-linearly with data duplication: sequences that appear more frequently in training data are reproduced more faithfully. For novels, the Gutenberg source text appears once. But for canonical novels, specific passages are *duplicated* across thousands of secondary sources: study guides that excerpt key quotations, academic papers that discuss them, Wikipedia articles that summarise plots with embedded excerpts, blog posts, teaching materials, and social media.

A famous passage—the fog opening of *Bleak House*, the pier glass passage of *Middlemarch*—plausibly appears not only in the Gutenberg text but across many such sources. We have not counted these occurrences, but the existence of study guides, Wikipedia articles, and academic papers for canonical novels, and their absence for obscure ones, makes the asymmetry likely. A passage from Chapter 12 of *Hester* appears on Gutenberg and, we conjecture, in few if any other sources. Our hypothesis: the LLM's per-novel quote fidelity tracks the volume of this **analytical reinforcement**, not the presence of the source text.

A.2 Data Sources and Composite Score

We constructed a composite attention score from three proxies, each capturing a different facet of a novel's analytical footprint:

1. **Goodreads ratings count.** Source: Goodreads.com, accessed March 2026. This measures present-day readership. Ratings counts range from 548 (*Hester*) to

[1] https://www.gutenberg.org. Gutenberg texts are a standard component of web-scraped training corpora such as C4 and The Pile. We assume all 15 novels are present in the training data of the model used (Claude, Anthropic), though we cannot verify this for any proprietary model.

261,000 (*David Copperfield*). Log-normalised to $[0, 100]$ to reduce the influence of extreme popularity: $\text{score} = 100 \times (\log r - \log r_{\min})/(\log r_{\max} - \log r_{\min})$ where r is the ratings count.

2. **Wikipedia article word count.** Source: English Wikipedia, accessed March 2026. Word count of the novel's Wikipedia article, measured by selecting all text content. This measures encyclopaedic and analytical attention: longer articles contain more plot summary, critical discussion, and quoted excerpts. Ranges from ~850 (*Miss Marjoribanks*) to ~19,000 (*David Copperfield*). Linearly normalised to $[0, 100]$.

3. **Study guide presence.** Binary indicator: does the novel have a SparkNotes, LitCharts, CliffsNotes, or Shmoop study guide? Source: manual check of each platform, March 2026. Study guides are particularly relevant because they systematically excerpt and re-quote key passages, producing exactly the kind of passage-level duplication that Carlini et al. (2023)'s framework predicts will strengthen memorisation. Coded as 33.3 (present) or 0 (absent).

The composite score is the unweighted mean of the three normalised components, producing a scale of 0–100. We acknowledge this scoring is ad hoc: the choice of proxies, normalisation methods, and equal weighting is not theoretically derived. Different reasonable choices would produce different numbers. We report it as one defensible operationalisation, not as a validated instrument.

A.3 Results

Table 3 presents the full data.

Overall correlation. Pearson $r = 0.72$, Spearman $\rho = 0.70$ ($n = 15$). Excluding *North and South* (discussed below): Pearson $r = 0.89$, Spearman $\rho = 0.90$ ($n = 14$).

A.4 The North and South Outlier

North and South has an attention score of 79.2—comparable to *Middlemarch* and *Bleak House*—but a nop verification rate of only 2%. Its Goodreads count (182,000) is almost certainly inflated by the 2004 BBC television adaptation, which generated a large modern fandom that discusses the adaptation's casting, romance, and divergences from the novel—not the novel's prose passage by passage.

This suggests that the composite attention score, which treats Goodreads popularity as equivalent to Wikipedia coverage and study guide presence, conflates *readership*

Table 3: Critical attention metrics and ungrounded (no-passages) quote verification rates. Data sources: Goodreads.com, English Wikipedia, SparkNotes/LitCharts (all accessed March 2026). Nop% is the percentage of LLM-generated quotes verified against the full enriched passage corpus using 5-tier matching (≥60% = verified).

Novel	*Author*	*Year*	*GR ratings*	*Wiki words*	*SG*	*Attn.*	*Nop%*
Middlemarch	Eliot	1871	181,000	8,750	Y	79.2	57
Bleak House	Dickens	1853	98,000	8,750	Y	75.9	53
David Copperfield	Dickens	1850	261,000	19,000	Y	100.0	44
Hard Times	Dickens	1854	75,000	6,750	Y	70.8	39
Passage to India	Forster	1924	73,000	4,350	Y	66.2	39
Mill on the Floss	Eliot	1860	59,000	3,000	Y	62.6	37
Cranford	Gaskell	1853	48,000	2,900	N	27.9	33
Daniel Deronda	Eliot	1876	27,000	3,500	Y	59.3	30
Our Mutual Friend	Dickens	1865	32,000	8,750	Y	69.8	29
New Grub Street	Gissing	1891	6,300	3,200	N	17.5	20
The Odd Women	Gissing	1893	5,500	1,200	N	13.1	6
Miss Marjoribanks	Oliphant	1866	2,600	850	N	8.4	3
North and South	Gaskell	1855	182,000	8,750	Y	79.2	2
Hester	Oliphant	1883	548	1,300	N	0.8	0
No Name	Collins	1862	7,900	3,200	N	18.7	0

with *analytical engagement*. A novel can be widely read (or widely watched) without generating the passage-level discussion that reinforces specific textual fragments in training data. *Cranford* (also adapted by the BBC, 2007) shows a milder version of this effect but remains closer to the trend line.

We report the outlier-excluded correlation ($r = 0.89$) alongside the full-sample correlation ($r = 0.72$), but we do not rely on the exclusion. The more important evidence is the within-author analysis below.

A.5 Within-Author Gradients

The most damaging objection to the analytical reinforcement hypothesis is that the correlation might reflect *author effects*: perhaps Dickens's prose is inherently more memorable (whether by LLMs, by people or by both) than Oliphant's, regardless of secondary discussion. We address this by examining within-author gradients.

Dickens (4 novels): Bleak House (53%) > David Copperfield (44%) > Hard Times (39%) > Our Mutual Friend (29%). This ordering is consistent with the attention scores (75.9, 100.0, 70.8, 69.8), though David Copperfield's high attention score (driven by its large

Table 4: Within-author nop verification gradients. We assume that novels by the same author share significant stylistic properties; author effects are mitigated.

Author	*Novel*	*Attn.*	*Nop%*
Dickens	Bleak House	75.9	53
	David Copperfield	100.0	44
	Hard Times	70.8	39
	Our Mutual Friend	69.8	29
Eliot	Middlemarch	79.2	57
	Mill on the Floss	62.6	37
	Daniel Deronda	59.3	30
Gaskell	Cranford	27.9	33
	North and South	79.2	2
Gissing	New Grub Street	17.5	20
	The Odd Women	13.1	6
Oliphant	Miss Marjoribanks	8.4	3
	Hester	0.8	0

Goodreads readership and long Wikipedia article) does not translate into the highest nop rate. The gradient exists within a single author's prose style, which is evidence against the author-effect confound, though with only four points the pattern could be coincidental.

Eliot (3 novels): Middlemarch (57%) > Mill on the Floss (37%) > Daniel Deronda (30%). The ordering tracks the attention scores (79.2, 62.6, 59.3). Middlemarch has the highest Goodreads count (181,000 vs. 59,000 and 27,000) and a much larger Wikipedia article than Daniel Deronda.

Gaskell (2 novels): Cranford (33%) ≫ North and South (2%). This is the most striking within-author contrast and the one we understand least well. North and South has more study guide coverage (GradeSaver, LitCharts, Course Hero, SuperSummary) and a higher attention score (79.2 vs. 27.9) than Cranford, yet Cranford dramatically outperforms. North and South's modern Goodreads popularity is inflated by the 2004 BBC adaptation, which may generate discussion of the *adaptation* without reinforcing the *novel's* specific passages. But we cannot rule out other explanations: Cranford's shorter length may make it more quotable per passage, or its comic tone may produce more distinctive (and therefore more memorisable) phrasing. This pair is genuinely anomalous and resists the simple analytical-reinforcement explanation.

Gissing (2 novels): New Grub Street (20%) > The Odd Women (6%). New Grub Street has a slightly higher attention score (17.5 vs. 13.1), driven by marginally more Goodreads ratings (6,300 vs. 5,500) and a longer Wikipedia article (3,200 vs. 1,200 words). Neither novel has a study guide.

Oliphant (2 novels): Miss Marjoribanks (3%) ≈ Hester (0%). Both are near zero, with the lowest attention scores in the corpus (8.4 and 0.8). Hester has only 548 Goodreads ratings, the fewest of any novel in the sample.

Within-author gradients exist for Dickens (4 novels), Eliot (3 novels), and Gissing (2 novels), and in each case the ranking is consistent with the attention scores. The Gaskell pair is anomalous (A.4), and the Oliphant pair is uninformative (both near zero). **The within-author evidence is the strongest argument against the author-effect confound**, because it holds prose style approximately constant. However, even within an author, novels differ in ways beyond critical prominence—length, subject matter, narrative voice—and we

cannot isolate analytical reinforcement as the sole explanation.

A.6 Passage-Level Evidence

We can also examine which *individual passages* the LLM reproduces most faithfully from memory. Across all *Bleak House* no-passages runs, the most reliably quoted passages are the fog opening of Chapter 1 (96% verified across 51 quote attempts) and Jo's death in Chapter 47 ("Dead, your Majesty…"; 92% across 78 attempts). These are among the novel's most celebrated passages, widely excerpted in study guides and anthologies.

This is **consistent with** the analytical reinforcement hypothesis: if passage-level duplication in secondary sources drives memorisation, the most-duplicated passages should be the most faithfully reproduced. However, it does not **test** the hypothesis directly. These passages also have distinctively rhythmic prose—the fog passage uses sustained anaphora, the death passage uses rhetorical repetition—and occupy structurally prominent positions in the novel (opening, climactic death). Either stylistic memorability or structural prominence could explain the result without recourse to secondary-source duplication.

Disentangling these explanations would require measuring passage-level frequency across the web—counting how many times each passage is quoted in secondary sources—and testing whether that frequency predicts nop fidelity after controlling for prose distinctiveness and structural position. We have not done this.

A.7 Objections and Qualifications

We subjected the thesis to systematic critique. The principal objections, and our responses:

Training data opacity. We do not know the exact composition of the model's training data. The attention score is a proxy for likely duplication in web-scraped corpora, not a measurement of actual duplication. The hypothesis is consistent with Carlini et al. (2023)'s findings but cannot be verified for proprietary models.

Ad hoc scoring. The composite score is one defensible operationalisation among several. Different weighting schemes or different proxies would produce different correlations. We have not conducted a systematic sensitivity analysis. A more principled instrument would require validated cultural-prominence metrics that do not currently exist.

Small sample. $n = 15$ supports trend identification but not confident statistical inference. The correlation ($r = 0.72$) has a wide confidence interval. We rely on the within-author evidence (Table 4) as much as on the aggregate correlation.

Readership vs. analytical engagement. The North and South outlier demonstrates that popular readership does not entail the kind of passage-level analytical engagement that would reinforce specific fragments in training data. A more refined metric would distinguish passage-quoting activity (study guides, academic close readings) from generic discussion (reviews, adaptation commentary). We have not measured this distinction directly.

The Bayard parallel is a metaphor. Bayard (2007)'s argument concerns human cultural navigation; our findings concern statistical weight distributions in a language model. The parallel—both humans and LLMs navigate literary culture through graduated approximation—is evocative but the mechanisms differ entirely. We flag it as an analogy, not an identity. As Shanahan (2024) argues, care is needed in mapping LLM behaviour onto human cognitive categories.

A.8 Conclusions

Stated conservatively:

All 15 novels are available from Project Gutenberg and plausibly appear in LLM training data. Yet the LLM's ability to quote them from memory varies from 57% to 0%. This gradient correlates with a composite measure of critical-cultural prominence

($r = 0.72$, $n = 15$). Within-author gradients for Dickens (4 novels), Eliot (3 novels), and Gissing (2 novels) are consistent with the attention scores, though the Gaskell pair is anomalous and the sample sizes are too small for confident within-author inference.

We interpret this as a literary-critical manifestation of differential LLM memorisation: canonical novels, whose passages are duplicated across study guides, academic papers, and Wikipedia, are more faithfully retained than obscure ones. This interpretation is consistent with Carlini et al. (2023)'s framework but unverifiable without access to training data.

The parallel with Bayard (2007)'s argument about human non-reading is suggestive: both humans and LLMs produce more confident discourse about texts they have engaged with more thoroughly. The mechanisms are entirely different—cultural positioning for humans, statistical weight distributions for models—and the parallel is an analogy, not an equivalence.

BRANJE BREZ BRANJA: KAKO BRALNA STRATEGIJA, INTERPRETATIVNI OKVIR IN BESEDILNA UTEMELJENOST OBLIKUJEJO RAČUNALNIŠKO LITERARNO RAZPRAVO

Bralec, ki roman prebere v celoti, in bralec, ki se osredotoči le na ključne odlomke, zbereta različne dokaze za razpravo. Ko pa obe bralni strategiji izvedemo računalniško in ustvarimo večglasne literarne razprave o 15 romanih, so rezultati skoraj enaki: delež govora strokovnjakov je stabilen ne glede na strategijo. To razhajanje med bralnimi strategijami in zbliževanje razprav zagotavlja računalniške dokaze za Fishovo tezo, da pomen konstruira bralec.

Ključne besede: literarna interpretacija, teorija bralnega odziva, memorizacija jezikovnih modelov, besedilna utemeljenost, digitalna humanistika